



Who is the best person for this? Answered from what people actually did, with the receipt attached.

hyderabaddies

Carl Kho, Steven Yang

September 27, 2026

At Recruit Holdings, who does what is decided from memory, as AI dissolves the roles.

WHO	Managers deciding placement and promotion, and the people decided about.
WHAT	Will is asked, never measured; matching runs on intuition; the evaluated build the case.
WHY NOW	AI dissolves job categories; teams now form around who can do the task. 80.9% of employees in Japan already use AI at work
WHY CURRENT OPTIONS FALL SHORT	Skill tags end in intuition; the evidence stays unused: one-sided use is “really creepy”.



Sources: Recruit Management Solutions, Jan 2026, n=1,429 · interviews at the venue, Sep 25 to 27, 2026, roles only

The evidence already exists. What was missing is a page the employee can see too.

OBSERVED

“I create a presentation of all the things that even my boss didn’t see.”

a product leader, 30 to 40 reports

REVEALED

“In theory we can use it. Accountability, transparency.”

a head of HR


UNLOCK

“They want to maintain their culture, they want to stick with their own rules of evaluation.”

a former leader of 150 engineers

Interviews at the venue, Sep 25 to 27, 2026, roles only · the Japanese original of the second quote is in the notes

Pik answers who is best for this, where you ask, with the receipt attached.



WHO IT IS FOR

The team lead, the HR planner, and the person named.

HOW IT WORKS

- 1 Ask where you already are.
- 2 Pik names the person, receipt attached.
- 3 The person named sees the same page.

Demo, 90 s: [deck/pik_demo_90s.mp4](#) · Meet listener on Gemini 3.8 Live, Vertex AI · the company on screen is fictional

Workers do the work.


Managers decide from receipts.

Nobody is scored.

51.2%

of HR staff in Japan say placement needs rethinking. Recruit's own survey, 2,761 respondents.

After an internal move, 75% are still there at two years, versus 56% without one.



TAM = talent management software, global, 2025
SAM = Japan HRTech cloud FY2024, ¥150 per USD = \$0.92B
SOM = 45,586 employees × \$8 per seat × 12 months

Score tools grade in secret. Pik shows both sides the receipt.

People Skills	Pik No score. Same page.
Lattice, Kaonavi, SmartHR	Workday, Gloat, Eightfold, Glean Delve, retired 2024

Hard to copy: a score product must unwind itself.

`lattice.com · microsoft.com: Delve retired 2024-12-16, People Skills GA 2025 · eightfold.ai · glean.com · pa-consul.co.jp`

The model extracts and cites. Software checks the quote. A human decides.



Python 3 stdlib · HTML + d3 · google-genai, Vertex AI: Gemini 3.8 Flash + Live · sounddevice · slack_bolt · notion-client

Recruit first. Then every floor that asks who has done this.



Notes: Birkmeyer, NEJM 2003 · Deltek Clarity 2022 · Recruit Holdings annual report FY2024, “Skills-First Hiring”

**Because the future of your
work depends on who you .**

Appendix

A1 Data policy

A2 Is that a score?

A3 What if it is wrong?

A4 What is real in the demo

A5 Japanese and English

A6 Deployment and security

A7 Pilot plan and metrics

A8 Competitors in detail

A9 The interviews

A10 Inspiration, alt

What Pik reads, and what it never reads.

READ, UNDER AN ALLOWLIST AND A PURPOSE LOCK

Public Slack channels; shared documents; Will Can Must sheets; manager notes beside the person's own words; opted-in AI sessions.

NEVER READ

Direct messages and private channels. Counted from a manifest, never opened.

prototype/policy.json · one named decision per page, deleted when it is made · every record change is logged

A2 IS THAT A SCORE?

No. It is coverage of the ask by a person's own receipts.

THE ENGINE COMPUTES

coverage = words of the ask with a receipt from this person ÷ words
in the ask

NEVER

Sorted into a rank, or shown to the person as a number. Tiles are listed
A to Z.

prototype/engine.py, coverage(): "2 of 7 words in the ask have a receipt from this person", per ask only

A3 WHAT IF IT IS WRONG?

The person sees the same page and can contest a line before the decision.

BEFORE THE DECISION

A note box on every line; the note lands on the evaluator's page.

INSIDE THE PIPELINE

A claim whose quote is not in its source is dropped, and the drops are counted.

In the demo: her manager wrote "flexible on location"; she wrote "I want to work with the US product team"

A4 WHAT IS REAL IN THE DEMO


The company is fictional. The pipeline and the surfaces run in the repo.

FIXTURES


Kaede Works, 200 generated people: 34 Slack messages, 5 documents, 3 sheets, 3 notes, 5 sessions, 2 DMs never read.

RUNNING


The Gemini 3.8 Live listener, Gemini 3.8 Flash extraction, the Slack bot, the Notion watcher, the page.



Meet



Slack



Notion

prototype/data/manifest.json · the memo is templated · connectors proposed; the demo reads JSON shaped like exports

The listener reacts in English, Japanese and Korean. Receipts stay in their own language.

AN ASK, AS SPOKEN

誰が英語で会議をリードできる？

Who can lead a meeting in English?

RECEIPTS

Never translated: the person's own sentence, in the language they wrote it. Pik speaks only when addressed.

Gemini 3.8 Live maps the ask onto the company's own skill tags; the engine does the rest · [prototype/live.py](#)

Inside the customer's Google Cloud project. Nothing leaves.

WHERE IT RUNS

In the customer's Google Cloud project on Vertex AI: Gemini 3.8 Flash extracts, Gemini 3.8 Live listens.

WHAT LEAVES

Nothing. The pilot runs on one laptop from the customer's exports; every change to a person record is logged.

Sources are allowlisted per decision; a purpose lock names the decision the page exists for · `prototype/policy.json`

One HR team, one decision type, one real cycle. Three measures.

PLAN

One HR team in Japan, their rubric and exports;
announced to the people concerned first.

MEASURES, ON ONE REAL CYCLE

Dropped-claim share, contest rate, evaluator minutes
per decision.

Dropped-claim share = claims rejected by the validator ÷ claims proposed · contest rate = lines annotated ÷ lines shown

Who reads what, who scores, what the person sees.

	READS	PRODUCES	PERSON SEES
Lattice	typed reviews	scores	their own review
Microsoft	Office 365	skill tags; Delve retired 2024	tags to confirm
Workday, Gloat, Eightfold	profiles	ranked matches	their profile
Glean	documents, tickets	ranked experts, opaque	nothing
Kaonavi, Talent Palette, SmartHR	HR data	match degree	not stated
Pik	what people did	the receipt, no score	the same page

Five conversations, about three hours twenty minutes, Friday and Saturday.

WHO

A head of HR, with an interpreter, 72 minutes

A product leader, 30 to 40 people, 48 minutes

A recruiter, new-graduate global talent acquisition, 40 minutes

A mentor, formerly a leader of up to 150 engineers, 25 minutes

A software engineer on a job-board team, 14 minutes

RULE

Roles only, never names.

Read in the order they come to mind, or laid out flat with the source on every line.

FROM MEMORY



FROM RECEIPTS

#northwind-sync, 2026-05-06	#learning, 2026-07-15	#growth-analytics, 2026-04-18
#northwind-sync, 2026-05-27	#growth-analytics, 2026-06-11	#pricing, 2026-08-12
#all-hands-questions, 2026-06-25	Will Can Must sheet, 2026-03-15	#platform, 2026-04-22

“It is difficult to search from the big data, so basically we depending based on our human knowledge.”

a recruiter, on how the one insider who has done this is found today · interviews at the venue, Sep 2026