



**Who is the best person for this? Answered from what people actually did, with the receipt attached.**

hyderabaddies

Carl Kho, Steven Yang

September 27, 2026

# At Recruit Holdings, will has no metric, and the evidence that shows it goes unread.

WHO About 3,000 first-line managers deciding placement, and the people decided about.

WHAT Will is asked twice a year, never measured; the evaluated build their own case.

WHY NOW Per-task task forces; the answer has to live in the tools already in use.

80.9% of employees in Japan already use AI at work

WHY CURRENT OPTIONS FALL SHORT Skill tags end in intuition; one-sided use of the evidence is “really creepy”.



one task force, per task

Recruit Management Solutions, Jan 2026, n=1,429 · 3,000: a VP of HR at Recruit Holdings, in interview · roles only

# The evidence already exists. What was missing is a page the employee can see too.

OBSERVED

**“I create a presentation of all the things that even my boss didn’t see.”**

a product leader, 30 to 40 reports

REVEALED

**“In theory we can use it. Accountability, transparency.”**

a head of HR


UNLOCK

**“They want to maintain their culture, they want to stick with their own rules of evaluation.”**

a former leader of 150 engineers

Interviews at the venue, Sep 25 to 27, 2026, roles only · the Japanese original of the second quote is in the notes

# Pik answers who is best for this, where you ask, with the receipt attached.



WHO IT IS FOR

The team lead, the HR planner, and the person named.

WHAT CHANGES

The lead: a receipt where the task lives.

The planner: a filed conclusion, questions first.

The person named: the same page, first.

One engine behind every surface; it ranks nothing.

Demo, 90 s: [deck/pik\\_demo\\_90s.mp4](#) · Meet listener on Gemini 3.8 Live, Vertex AI · the company on screen is fictional

**Workers do the work.**

**Managers decide from receipts.**


**Nobody is scored.**

# 51.2%

of HR staff in Japan say placement needs rethinking. Recruit's own survey, 2,761 respondents.

After an internal move, 75% are still there at two years, versus 56% without one.

In interview, HR leadership named ¥200M for HR, training and matching systems, and ¥100M per exchange posting.



TAM = talent management software, global, 2025

SAM = Japan HRTech cloud FY2024, ¥150 per USD = \$0.92B

SOM =  $45,586 \times \$8 \times 12 = \$4.4\text{M a year}$ ; Japan only 12,709 = \$1.2M

Interview: ¥200M = \$1.3M, ¥100M = \$0.67M at ¥150; no period stated

Recruit 2023 · LinkedIn 2023 · Fortune BI 2025 · MIC 2025 · S&P · Lattice · VP of HR at Recruit, interview · not to scale

# Score tools grade in secret. Pik shows both sides the receipt.

|                           |                                                         |
|---------------------------|---------------------------------------------------------|
| People Skills             | <b>Pik</b><br>No score. Same page.                      |
| Lattice, Kaonavi, SmartHR | Workday, Gloat, Eightfold, Glean<br>Delve, retired 2024 |


SCORE → SAME PAGE

PROFILES → WHAT PEOPLE DID

Hard to copy: a score product must unwind itself.


`lattice.com · microsoft.com: Delve retired 2024-12-16, People Skills GA 2025 · eightfold.ai · glean.com · pa-consul.co.jp`

# The model extracts and cites. Software checks the quote. A human decides.



Python 3 stdlib · HTML + d3 · google-genai, Vertex AI: Gemini 3.8 Flash + Live · sounddevice · slack\_bolt · notion-client

# Next: a learned employee state, finding what keywords miss.



A VERBATIM QUOTE WITH ITS SOURCE

NO NUMBER ABOUT ANYONE

SAME PAGE; CONTESTS ARE LABELS

Trained inside the customer's Vertex AI project; evaluation on held-out past decisions; nothing leaves.

# Recruit first. Then every floor that asks who has done this.



Notes: Birkmeyer, NEJM 2003 · Deltek Clarity 2022 · Recruit Holdings annual report FY2024, “Skills-First Hiring”

**Same page for both sides. A receipt on every line.**

# Appendix

A1 Data policy

A2 Is that a score?

A3 What if it is wrong?

A4 What is real in the demo

A5 Japanese and English

A6 Deployment and security

A7 Pilot plan and metrics

A8 Competitors in detail

A9 The interviews

A10 Inspiration, alt

A11 The state model

# What Pik reads, and what it never reads.

READ, UNDER AN ALLOWLIST AND A PURPOSE LOCK

Public Slack channels; shared documents; Will Can Must sheets; manager notes beside the person's own words; opted-in AI sessions.

NEVER READ

Direct messages and private channels. Counted from a manifest, never opened.

prototype/policy.json · one named decision per page, deleted when it is made · every record change is logged

A2 IS THAT A SCORE?

**No. It is coverage of the ask by a person's own receipts.**

THE ENGINE COMPUTES

coverage = words of the ask with a receipt from this person ÷ words in the ask

NEVER

Sorted into a rank, or shown to the person as a number. Tiles are listed A to Z.

prototype/engine.py, coverage(): "2 of 7 words in the ask have a receipt from this person", per ask only

A3 WHAT IF IT IS WRONG?

# The person sees the same page and can contest a line before the decision.

BEFORE THE DECISION

A note box on every line; the note lands on the evaluator's page.

INSIDE THE PIPELINE

A claim whose quote is not in its source is dropped, and the drops are counted.

In the demo: her manager wrote "flexible on location"; she wrote "I want to work with the US product team"


# The company is fictional. The pipeline and the surfaces run in the repo.

## FIXTURES


Kaede Works, 200 generated people: 34 Slack messages, 5 documents, 3 sheets, 3 notes, 5 sessions, 2 DMs never read.

## RUNNING


The Gemini 3.8 Live listener, Gemini 3.8 Flash extraction, the Slack bot, the Notion watcher, the page.



Meet



Slack



Notion

prototype/data/manifest.json · the memo is templated · connectors proposed; the demo reads JSON shaped like exports

# The listener reacts in English, Japanese and Korean. Receipts stay in their own language.

AN ASK, AS SPOKEN

誰が英語で会議をリードできる？

Who can lead a meeting in English?

RECEIPTS

Never translated: the person's own sentence, in the language they wrote it. Pik speaks only when addressed.

Gemini 3.8 Live maps the ask onto the company's own skill tags; the engine does the rest · [prototype/live.py](#)

# Inside the customer's Google Cloud project. Nothing leaves.

## WHERE IT RUNS

In the customer's Google Cloud project on Vertex AI: Gemini 3.8 Flash extracts, Gemini 3.8 Live listens.

## WHAT LEAVES

Nothing. The pilot runs on one laptop from the customer's exports; every change to a person record is logged.

Sources are allowlisted per decision; a purpose lock names the decision the page exists for · `prototype/policy.json`

# One HR team, one decision type, one real cycle. Three measures.

## PLAN

One HR team in Japan, their rubric and exports;  
announced to the people concerned first.

## MEASURES, ON ONE REAL CYCLE

Dropped-claim share, contest rate, evaluator minutes  
per decision.

Dropped-claim share = claims rejected by the validator ÷ claims proposed · contest rate = lines annotated ÷ lines shown

# Who reads what, who scores, what the person sees.

|                                  | READS              | PRODUCES                       | PERSON SEES      |
|----------------------------------|--------------------|--------------------------------|------------------|
| Lattice                          | typed reviews      | scores                         | their own review |
| Microsoft                        | Office 365         | skill tags; Delve retired 2024 | tags to confirm  |
| Workday, Gloat, Eightfold        | profiles           | ranked matches                 | their profile    |
| Glean                            | documents, tickets | ranked experts, opaque         | nothing          |
| Kaonavi, Talent Palette, SmartHR | HR data            | match degree                   | not stated       |
| Pik                              | what people did    | the receipt, no score          | the same page    |

# Five conversations, about three hours twenty minutes, Friday and Saturday.

## WHO

A head of HR, with an interpreter, 72 minutes

A product leader, 30 to 40 people, 48 minutes

A recruiter, new-graduate global talent acquisition, 40 minutes

A mentor, formerly a leader of up to 150 engineers, 25 minutes

A software engineer on a job-board team, 14 minutes

## RULE

Roles only, never names.

# Read in the order they come to mind, or laid out flat with where each came from.

FROM MEMORY



FROM RECEIPTS

|                                                                              |                                                                             |                                                                           |
|------------------------------------------------------------------------------|-----------------------------------------------------------------------------|---------------------------------------------------------------------------|
| <div><div></div><div></div><div>#northwind-sync, 2026-05-06</div></div>      | <div><div></div><div></div><div>#learning, 2026-07-15</div></div>           | <div><div></div><div></div><div>#growth-analytics, 2026-04-18</div></div> |
| <div><div></div><div></div><div>#northwind-sync, 2026-05-27</div></div>      | <div><div></div><div></div><div>#growth-analytics, 2026-06-11</div></div>   | <div><div></div><div></div><div>#pricing, 2026-08-12</div></div>          |
| <div><div></div><div></div><div>#all-hands-questions, 2026-06-25</div></div> | <div><div></div><div></div><div>Will Can Must sheet, 2026-03-15</div></div> | <div><div></div><div></div><div>#platform, 2026-04-22</div></div>         |

“It is difficult to search from the big data, so basically we depending based on our human knowledge.”

a recruiter, on how the one insider who has done this is found today · interviews at the venue, Sep 2026

# A retrieval model over receipts, never a prediction about a person.

## WHAT IT LEARNS

What each person has done, wants and can take on, as a time-aware vector, from allowed receipts only.

## HOW IT IS TRAINED

Contrastive retrieval: asks paired with the receipts filed decisions used; contests and dropped claims as corrections.

Evaluated by recall of decision-used receipts on held-out cycles, contest rate, dropped-claim share.

Retrieval objective, not a prediction about a person; no protected attributes; the receipt is the unit; a human decides.