

01 TITLE



Who is the best person for this task? Answered from what people actually did, with the source attached.

hyderabaddies

September 27, 2026

At Recruit Holdings, every task needs its own task force, and nobody can see who fits.

WHO

Anyone staffing a task: team leads, HR, and every colleague who needs to loop someone in.

WHAT

Finding who has done this, and who wants to, for every new task.

WHY NOW

With AI, HR can ship code and engineers can run HR. Fixed teams give way to one task force per task.

80.9% of employees in Japan already use AI at work

WHY CURRENT OPTIONS
FALL SHORT

Directories and skill tags end in intuition, so the same few people get every task.



one task force, per task

The evidence already exists. What was missing is a page the employee can see too.



OBSERVED

“Every cycle I build my own deck of everything my boss didn’t see.”


a product leader who manages 40



REVEALED

“Team formation, as AI dissolves the roles. That is what we’re most troubled by.”

a head of HR




UNLOCK

“Managers pitch people from memory. Four or five days a quarter.”

a former engineering leader of 150

Interviews at the venue, Sep 25 to 27, 2026, roles only · the Japanese original of the second quote is in the notes

Pik finds who is best for each task, from the work your company's tools already hold.



WHO IT IS FOR

Anyone forming a task force: team leads, HR, and any colleague who needs to loop someone in.

WHAT CHANGES

Ask once, where you work. No directory search.

Teams form per task, then dissolve.

Work follows capacity, not the usual few.

Matched on what people did, want, and how they work together.

Demo, 90 s: [deck/pik_demo_90s.mp4](#) · Meet listener on Gemini 3.8 Live, Vertex AI · the company on screen is fictional

Job titles stop being walls.

Every task gets its own task force.


Everyone blooms where they fit.

51.2%

of HR staff in Japan say placement needs rethinking. Recruit's own survey, 2,761 respondents. [RECRUIT](#)

After an internal move, 75% are still there at two years, versus 56% without one. [in LinkedIn](#)

HR leadership named about \$1.3M spent on HR, training and matching systems. [🎙 Interview](#)



[Fortune Business Insights](#) TAM = talent management software, global, 2025

[Deloitte Tohmatsu MIC](#) SAM = Japan HR-tech cloud, FY2024

[S&P Global · Lattice](#) SOM = $45,586 \times \$8 \times 12 = \$4.4M$ a year

Yen converted at 150 per USD; interview figure has no period


Recruit 2023 · LinkedIn 2023 · Fortune BI 2025 · MIC 2025 · S&P · Lattice · VP of HR at Recruit, interview · not to scale

Others profile people. Pik computes a confidence per task, and shows why.

GRADE OF A PERSON → CONFIDENCE PER TASK	People Skills	Pik Confidence per task. Hard and soft skills, culture. Same page.
	Lattice, Kaonavi, SmartHR	Workday, Gloat, Eightfold, Glean Delve, retired 2024
	TYPED PROFILES →	WHAT PEOPLE DID


Hard to copy: it reads soft skills and culture from how people actually work, not from a profile.

Every signal is checked against its source, then weighed per task. A human decides.



Python · Vertex AI: Gemini 3.8 Flash; 3.8 Live joins Meet · Slack Bolt · Notion API · retriever opt-in: PIK_RETRIEVER_URL

Built: our own multilingual retriever. Next: more layers of capture.



CAPTURE IS OPT-IN

DMS NEVER READ

THE PERSON SEES THE SAME PAGE

Today: a fine-tuned retriever that answers Japanese and English asks (finetune/) · Next: opt-in OCR, VLM, connectors

The fine-tune, on one ask.

INPUT

“Who offered to facilitate the look-back session about A/B tests on what we charge?”

BEFORE · KEYWORD ENGINE, THE DEMO TODAY

matches “tests” and “what”

Yui Sato, for “Q2 pricing test: results and what I’d change”

Doc · doc-yui-2

then 2 people with nothing on file

asks back: “Nothing on file mentions about, charge, facilitate, look-back.”

AFTER · FINE-TUNED RETRIEVER

Five of five top results are right.

“Volunteered to run the pricing experiments retro this quarter so the team lead could focus on hiring.”

Jun Murakami · Slack #design-crit, 2026-07-04 (gr-044-0)


and 4 more people who said the same

WHAT BETTER MEANS

The right evidence in the top five: 51% → 85% on 364 held-out English asks; 0% → 70% on asks that share no word with it.

A held-out ask; fictional fixtures · output of finetune/ask.py · finetune/results/metrics.json

Fine-tuned, the right evidence is in the top five 85% of the time, up from 51%.



Held-out English asks (n=364) and Japanese asks (n=28), fictional fixtures, trained locally 2026-09-27; finetune/results/metrics.json

Recruit first. Then every floor that asks who has done this.



Notes: Birkmeyer, NEJM 2003 · Deltek Clarity 2022 · Recruit Holdings annual report FY2024, “Skills-First Hiring”

**Because the future of work
depends on who you **Pik**.**

Appendix

A1 Data policy

A2 Is that a score?

A3 What if it is wrong?

A4 What is real in the demo

A5 Japanese and English

A6 Deployment and security

A7 Pilot plan and metrics

A8 Competitors in detail

A9 The interviews

A10 Inspiration, alt

A11 The state model

What Pik reads, and what it never reads.

READ, UNDER AN ALLOWLIST AND A PURPOSE LOCK

Public Slack channels; Google Drive documents; Will Can Must sheets; manager notes beside the person's own words; opted-in AI sessions.

NEVER READ

Direct messages and private channels. Counted from a manifest, never opened.

prototype/policy.json · one named decision per page, deleted when it is made · every record change is logged

A2 IS THAT A SCORE?

No. It is a confidence for one task, from a person's own evidence.

THE ENGINE COMPUTES

coverage = words of the ask with evidence from this person ÷ words in the ask

NEVER

Sorted into a rank, or shown to the person as a number. Tiles are listed A to Z.

prototype/engine.py, coverage(): "2 of 7 words in the ask have evidence from this person", per ask only

A3 WHAT IF IT IS WRONG?

The person sees the same page and can contest a line before the decision.

BEFORE THE DECISION

A note box on every line; the note lands on the evaluator's page.

INSIDE THE PIPELINE

A claim whose quote is not in its source is dropped, and the drops are counted.

In the demo: her manager wrote "flexible on location"; she wrote "I want to work with the US product team"


The company is fictional. The pipeline and the surfaces run in the repo.

FIXTURES


Kaede Works, 200 generated people: 34 Slack messages, 5 documents, 3 sheets, 3 notes, 5 sessions, 2 DMs never read.

RUNNING


The Gemini 3.8 Live listener, Gemini 3.8 Flash extraction, the Slack bot, the Notion watcher, the page.



Meet



Slack



Notion

prototype/data/manifest.json · the memo is templated · connectors proposed; the demo reads JSON shaped like exports

The listener reacts in English, Japanese and Korean. Evidence stays in its own language.

AN ASK, AS SPOKEN

誰が英語で会議をリードできる？

Who can lead a meeting in English?

EVIDENCE

Never translated: the person's own sentence, in the language they wrote it. Pik speaks only when addressed.

Gemini 3.8 Live maps the ask onto the company's own skill tags; the engine does the rest · [prototype/live.py](#)

Inside the customer's Google Cloud project, or on an air-gapped server. Nothing leaves.

WHERE IT RUNS

The customer's Google Cloud project on Vertex AI, or an air-gapped server on their premises with our fine-tuned retriever.

WHAT LEAVES

Nothing. The pilot runs on one laptop from the customer's exports; every change to a person record is logged.

Sources are allowlisted per decision; a purpose lock names the decision the page exists for · `prototype/policy.json`

One HR team, one decision type, twelve weeks. Three measures.

PLAN

One HR team in Japan, their rubric and exports.
Weeks 1–2 set up, 3–10 one real cycle, 11–12 review.

MEASURES, ON ONE REAL CYCLE

Time to staff a task, before and with Pik; contest rate; dropped-claim share.

Dropped-claim share = claims rejected by the validator ÷ claims proposed · contest rate = lines annotated ÷ lines shown

Who reads what, who scores, what the person sees.

	READS	PRODUCES	PERSON SEES
Lattice	typed reviews	scores	their own review
Microsoft	Office 365	skill tags; Delve retired 2024	tags to confirm
Workday, Gloat, Eightfold	profiles	ranked matches	their profile
Glean	documents, tickets	ranked experts, opaque	nothing
Kaonavi, Talent Palette, SmartHR	HR data	match degree	not stated
Pik	what people did	confidence per task, sourced	the same page

Five conversations, about three hours twenty minutes, Friday and Saturday.

WHO

A head of HR, with an interpreter, 72 minutes

A product leader, 30 to 40 people, 48 minutes

A recruiter, new-graduate global talent acquisition, 40 minutes

A mentor, formerly a leader of up to 150 engineers, 25 minutes

A software engineer on a job-board team, 14 minutes

RULE

Roles only, never names.

Recorded at the venue, Sep 25 to 27, 2026 · figures are the speakers' own past organisations, not Recruit's

Read in the order they come to mind, or laid out flat with where each came from.

FROM MEMORY



FROM EVIDENCE

 #northwind-sync, 2026-05-06	 #learning, 2026-07-15	 #growth-analytics, 2026-04-18
 #northwind-sync, 2026-05-27	 #growth-analytics, 2026-06-11	 #pricing, 2026-08-12
 #all-hands-questions, 2026-06-25	 Will Can Must sheet, 2026-03-15	 #platform, 2026-04-22

“It is difficult to search from the big data, so basically we depending based on our human knowledge.”

a recruiter, on how the one insider who has done this is found today · interviews at the venue, Sep 2026

A retrieval model over evidence, never a prediction about a person.

WHAT IT LEARNS

What each person has done, wants and can take on, as a time-aware vector, from allowed evidence only.

HOW IT IS TRAINED

Contrastive retrieval: asks paired with the evidence filed decisions used; contests and dropped claims as corrections.

Evaluated by recall of decision-used evidence on held-out cycles, contest rate, dropped-claim share.

Retrieval, not a prediction about a person · no protected attributes · a human decides